

Abalone Age Prediction

Assignment -3

| | |
|---------------------|--|
| Assignment Date | 03 October 2022 |
| Team ID | PNT2022TMID27812 |
| Project Name | Smart Lender-Application Credibility Prediction for loan Approval |
| Student Name | Mugeshraja.M |
| Student Roll Number | 311519104702 |
| Maximum Marks | 2 Marks |

Question-1.Download dataset

Solution:

| 1 | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings | | |
|----|----------|--------|--------------|----------------|----------------|--------------|-------|--|--|
| 2 | 0.365 | 0.095 | 0.514 | 0.2245 | 0.101 | 0.15 | 15 | | |
| 3 | 0.265 | 0.09 | 0.2255 | 0.0995 | 0.0485 | 0.07 | 7 | | |
| 4 | 0.42 | 0.135 | 0.677 | 0.2565 | 0.1415 | 0.21 | 9 | | |
| 5 | 0.365 | 0.125 | 0.516 | 0.2155 | 0.114 | 0.155 | 10 | | |
| 6 | 0.255 | 0.08 | 0.205 | 0.0895 | 0.0395 | 0.055 | 7 | | |
| 7 | 0.3 | 0.095 | 0.3515 | 0.141 | 0.0775 | 0.12 | 8 | | |
| 8 | 0.415 | 0.15 | 0.7775 | 0.237 | 0.1415 | 0.33 | 20 | | |
| 9 | 0.425 | 0.125 | 0.768 | 0.294 | 0.1495 | 0.26 | 16 | | |
| 10 | 0.37 | 0.125 | 0.5095 | 0.2165 | 0.1125 | 0.165 | 9 | | |
| 11 | 0.44 | 0.15 | 0.8945 | 0.3145 | 0.151 | 0.32 | 19 | | |
| 12 | 0.38 | 0.14 | 0.6065 | 0.194 | 0.1475 | 0.21 | 14 | | |
| 13 | 0.35 | 0.11 | 0.406 | 0.1675 | 0.081 | 0.135 | 10 | | |
| 14 | 0.38 | 0.135 | 0.5415 | 0.2175 | 0.095 | 0.19 | 11 | | |
| 15 | 0.405 | 0.145 | 0.6845 | 0.2725 | 0.171 | 0.205 | 10 | | |
| 16 | 0.355 | 0.1 | 0.4755 | 0.1675 | 0.0805 | 0.185 | 10 | | |
| 17 | 0.4 | 0.13 | 0.6645 | 0.258 | 0.133 | 0.24 | 12 | | |
| 18 | 0.28 | 0.085 | 0.2905 | 0.095 | 0.0395 | 0.115 | 7 | | |
| 19 | 0.34 | 0.1 | 0.451 | 0.188 | 0.087 | 0.13 | 10 | | |
| 20 | 0.295 | 0.08 | 0.2555 | 0.097 | 0.043 | 0.1 | 7 | | |
| 21 | 0.32 | 0.1 | 0.381 | 0.1705 | 0.075 | 0.115 | 9 | | |
| 22 | 0.28 | 0.095 | 0.2455 | 0.0955 | 0.062 | 0.075 | 11 | | |
| 23 | 0.275 | 0.1 | 0.2255 | 0.08 | 0.049 | 0.085 | 10 | | |
| 24 | 0.44 | 0.155 | 0.9395 | 0.4275 | 0.214 | 0.27 | 12 | | |
| 25 | 0.415 | 0.135 | 0.7635 | 0.318 | 0.21 | 0.2 | 9 | | |

Question-2.Load the dataset

Solution:

```
import numpy as np
```

```
import pandas as pd
```

```
import seaborn as sns
```

```
import matplotlib.pyplot as plt
```

```
import sklearn
```

```
data = pd.read_csv(r"abalone.csv")
```

```
data.head
```

```
<bound method NDFrame.head of
0      M  0.455  0.365  0.095  0.5140  0.2245
1      M  0.350  0.265  0.090  0.2255  0.0995
2      F  0.530  0.420  0.135  0.6770  0.2565
3      M  0.440  0.365  0.125  0.5160  0.2155
4      I  0.330  0.255  0.080  0.2050  0.0895
...
4172   F  0.565  0.450  0.165  0.8870  0.3700
4173   M  0.590  0.440  0.135  0.9660  0.4390
4174   M  0.600  0.475  0.205  1.1760  0.5255
4175   F  0.625  0.485  0.150  1.0945  0.5310
4176   M  0.710  0.555  0.195  1.9485  0.9455


Viscera weight  Shell weight  Rings
0              0.1010      0.1500    15
1              0.0485      0.0700     7
2              0.1415      0.2100     9
3              0.1140      0.1550    10
4              0.0395      0.0550     7
...
4172           0.2390      0.2490    11
4173           0.2145      0.2605    10
4174           0.2875      0.3080     9
4175           0.2610      0.2960    10
4176           0.3765      0.4950    12
```

Question-3. Perform Below Visualizations.


3.1 Univariate Analysis

Solution:

`sns.displot(data[' Diameter'])`

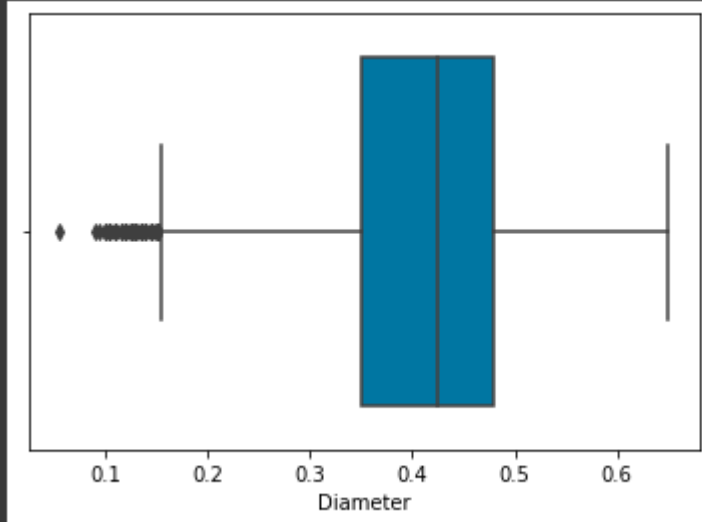


`sns.histplot(data[' Diameter '])`



```
sns.boxplot(x = data[' Diameter '])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4731587050>
```




3.2 Bivariate Analysis

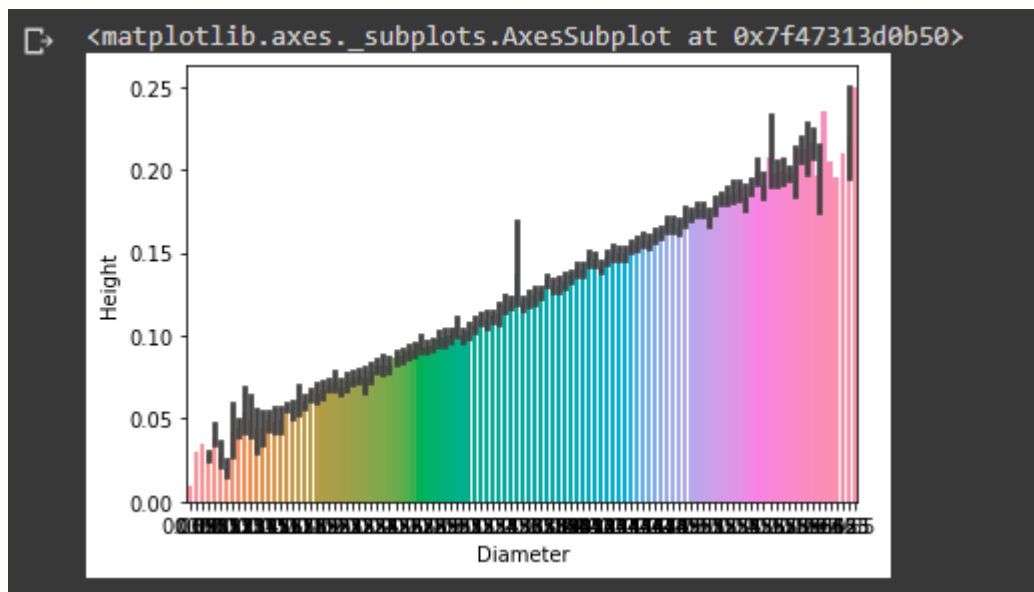
Solution:

```
sns.lineplot(data = data, x = 'Diameter', y = 'Rings')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f47314ec350>
```




```
sns.barplot(data = data, x = 'Diameter', y = 'Height')
```



```
plt.figure(figsize=(20,20))
```


```
sns.barplot(data = data, x = 'Diameter', y = 'Shell weight')
```




3.3 Multivariate Analysis

Solution:

```
sns.scatterplot(data = data, x = 'Shell weight', y = 'Rings', hue = 'Length')
```



sns.pairplot(data)



Question-4. Perform descriptive statistics on the dataset.

Solution:

```
data.mean(numeric_only = True)
```

```
Length      0.523992
Diameter    0.407881
Height      0.139516
Whole weight 0.828742
Shucked weight 0.359367
Viscera weight 0.180594
Shell weight 0.238831
Rings       9.933684
dtype: float64
```

```
data.median(numeric_only = True)
```

```
Length      0.5450
Diameter    0.4250
Height      0.1400
Whole weight 0.7995
Shucked weight 0.3360
Viscera weight 0.1710
Shell weight 0.2340
Rings       9.0000
dtype: float64
```

```
data['Whole weight'].mode()
```

```
0 0.2225 dtype: float64
```

```
data['Length'].mode()
```

```
0 0.550 1 0.625 dtype: float64
```

```
data['Rings'].unique()
```

```
array([15, 7, 9, 10, 8, 20, 16, 19, 14, 11, 12, 18, 13, 5, 4, 6, 21,
       17, 22, 1, 3, 26, 23, 29, 2, 27, 25, 24])
```

```
data.std(numeric_only=True)
```

```
Length      0.120093
Diameter    0.099240
Height      0.041827
Whole weight 0.490389
Shucked weight 0.221963
Viscera weight 0.109614
Shell weight 0.139203
Rings       3.224169
dtype: float64
```


`data.describe()`

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|-------|-------------|-------------|-------------|--------------|----------------|----------------|--------------|-------------|
| count | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 | 4177.000000 |
| mean | 0.523992 | 0.407881 | 0.139516 | 0.828742 | 0.359367 | 0.180594 | 0.238831 | 9.933684 |
| std | 0.120093 | 0.099240 | 0.041827 | 0.490389 | 0.221963 | 0.109614 | 0.139203 | 3.224169 |
| min | 0.075000 | 0.055000 | 0.000000 | 0.002000 | 0.001000 | 0.000500 | 0.001500 | 1.000000 |
| 25% | 0.450000 | 0.350000 | 0.115000 | 0.441500 | 0.186000 | 0.093500 | 0.130000 | 8.000000 |
| 50% | 0.545000 | 0.425000 | 0.140000 | 0.799500 | 0.336000 | 0.171000 | 0.234000 | 9.000000 |
| 75% | 0.615000 | 0.480000 | 0.165000 | 1.153000 | 0.502000 | 0.253000 | 0.329000 | 11.000000 |
| max | 0.815000 | 0.650000 | 1.130000 | 2.825500 | 1.488000 | 0.760000 | 1.005000 | 29.000000 |

`data['Whole weight'].value_counts()`

```
0.2225    8
1.1345    7
0.9700    7
0.4775    7
0.1960    7
..
0.0475    1
1.8930    1
1.8725    1
2.1055    1
1.9485    1
Name: Whole weight, Length: 2429, dtype: int64
```

Question-5.Handle the Missing values.

Solution:

data.isnull().any()

```
Sex          False
Length       False
Diameter     False
Height       False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Rings        False
dtype: bool
```

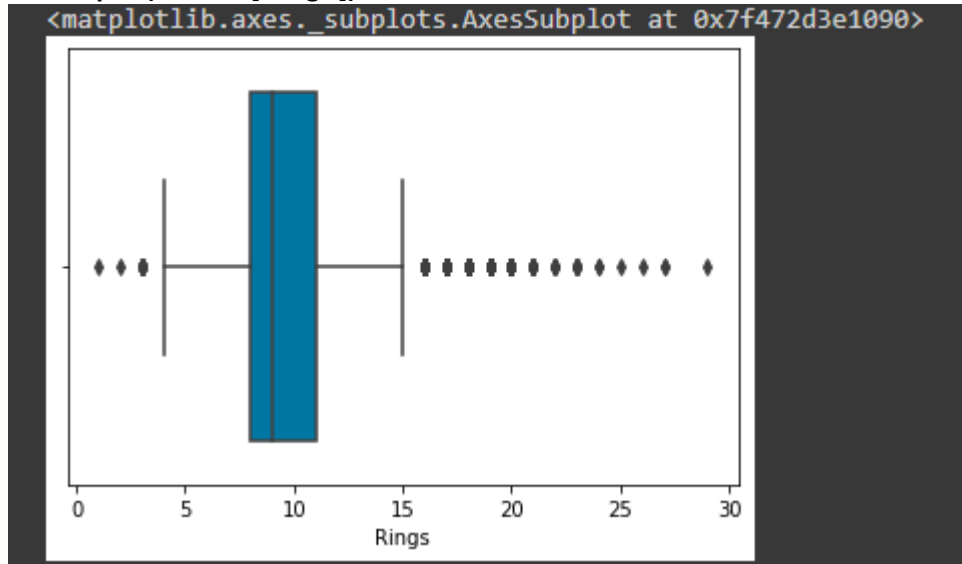
data.isnull().sum()

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings        0
dtype: int64
```

Question-6. Find the outliers and replace the outliers

Solution:

```
sns.boxplot(x = data['Rings'])
```



```
q = data.quantile([0.75,0.25])
```

q

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings |
|------|--------|----------|--------|--------------|----------------|----------------|--------------|-------|
| 0.75 | 0.615 | 0.48 | 0.165 | 1.1530 | 0.502 | 0.2530 | 0.329 | 11.0 |
| 0.25 | 0.450 | 0.35 | 0.115 | 0.4415 | 0.186 | 0.0935 | 0.130 | 8.0 |

```
iqr = q.iloc[0] - q.iloc[1]
```

iqr

```
Length      0.1650
Diameter    0.1300
Height      0.0500
Whole weight 0.7115
Shucked weight 0.3160
Viscera weight 0.1595
Shell weight 0.1990
Rings       3.0000
dtype: float64
```

```
u = q.iloc[0] + (1.5*iqr)
```

```
u
```

```
Length      0.86250
Diameter    0.67500
Height      0.24000
Whole weight 2.22025
Shucked weight 0.97600
Viscera weight 0.49225
Shell weight 0.62750
Rings       15.50000
dtype: float64
```


```
l = q.iloc[1] - (1.5*iqr)
```

```
l
```

```
Length      0.20250
Diameter    0.15500
Height      0.04000
Whole weight -0.62575
Shucked weight -0.28800
Viscera weight -0.14575
Shell weight -0.16850
Rings        3.50000
dtype: float64
```

```
data['Rings'] = np.where(np.logical_or(data['Rings']>15, data['Rings']<4), 9, data['Rings'])
sns.boxplot(x = data['Rings'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f472d26c310>
```



Question-7. Check for Categorical columns and perform encoding

Solution:

```
data['Age'] = data['Rings'] + 1.5  
data['Sex'].value_counts()
```

```
M    1528  
I    1342  
F    1307  
Name: Sex, dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder  
le = LabelEncoder()  
data = pd.get_dummies(data)  
data.head()
```

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Rings | Age | Sex_F | Sex_I | Sex_M |
|---|--------|----------|--------|--------------|----------------|----------------|--------------|-------|------|-------|-------|-------|
| 0 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 15 | 16.5 | 0 | 0 | 1 |
| 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 7 | 8.5 | 0 | 0 | 1 |
| 2 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 9 | 10.5 | 1 | 0 | 0 |
| 3 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 10 | 11.5 | 0 | 0 | 1 |
| 4 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 7 | 8.5 | 0 | 1 | 0 |

Question-8. Split the data into dependent and independent variables split the data in X and Y

Solution:

```
y = data['Age']  
data = data.drop(['Rings','Age'], axis=1)  
x = data  
data.head
```

| | Length | Diameter | Height | Whole weight | Shucked weight | Viscera weight | Shell weight | Sex_F | Sex_I | Sex_M |
|---|--------|----------|--------|--------------|----------------|----------------|--------------|-------|-------|-------|
| 0 | 0.455 | 0.365 | 0.095 | 0.5140 | 0.2245 | 0.1010 | 0.150 | 0 | 0 | 1 |
| 1 | 0.350 | 0.265 | 0.090 | 0.2255 | 0.0995 | 0.0485 | 0.070 | 0 | 0 | 1 |
| 2 | 0.530 | 0.420 | 0.135 | 0.6770 | 0.2565 | 0.1415 | 0.210 | 1 | 0 | 0 |
| 3 | 0.440 | 0.365 | 0.125 | 0.5160 | 0.2155 | 0.1140 | 0.155 | 0 | 0 | 1 |
| 4 | 0.330 | 0.255 | 0.080 | 0.2050 | 0.0895 | 0.0395 | 0.055 | 0 | 1 | 0 |

y

```
0      16.5
1       8.5
2     10.5
3     11.5
4       8.5
...
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5
Name: Age, Length: 4177, dtype: float64
```

Question-9.Scale the independent variables

Solution:

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled = sc.fit_transform(x)
x_scaled
```

```
array([[ -0.57455813, -0.43214879, -1.06442415, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-1.44898585, -1.439929  , -1.18397831, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.05003309,  0.12213032, -0.10799087, ...,  1.48184628,
        -0.68801788, -0.75948762],
       ...,
       [ 0.6329849  ,  0.67640943,  1.56576738, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.84118198,  0.77718745,  0.25067161, ...,  1.48184628,
        -0.68801788, -0.75948762],
       [ 1.54905203,  1.48263359,  1.32665906, ..., -0.67483383,
        -0.68801788,  1.31667716]])
```

Question-10.Split data into Training and Testing

Solution:

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.3, random_state = 0)
```

x_train

```
array([[ 0.79954256,  1.0291325 ,  0.84844242, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-1.49062526, -1.54070702, -1.30353247, ..., -0.67483383,
         1.45345059, -0.75948762],
       [-1.24078877, -1.33915098, -1.06442415, ..., -0.67483383,
         1.45345059, -0.75948762],
       ...,
       [ 0.59134549,  0.42446438,  0.13111745, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.84118198,  0.82757646,  0.6093341 , ...,  1.48184628,
        -0.68801788, -0.75948762],
       [-0.94931287, -0.83526087, -0.70576167, ..., -0.67483383,
         1.45345059, -0.75948762]])
```

x_train.shape

```
(2923, 10)
```

x_test

```
array([[ 0.21659075,  0.17251933,  0.37022577, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-0.1998034 , -0.07942572, -0.46665335, ..., -0.67483383,
         1.45345059, -0.75948762],
       [ 0.79954256,  0.72679844,  0.37022577, ..., -0.67483383,
        -0.68801788,  1.31667716],
       ...,
       [ 0.92446081,  0.87796547, -2.97729071, ...,  1.48184628,
        -0.68801788, -0.75948762],
       [ 1.13265788,  0.97874349,  1.44621322, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.79954256,  0.77718745,  0.72888826, ..., -0.67483383,
         1.45345059, -0.75948762]])
```

x_test.shape

```
(1254, 10)
```

y_train

```
1376    11.5
1225     6.5
2722     8.5
3387    10.5
2773    12.5
...
1033    11.5
3264    13.5
1653    11.5
2607    10.5
2732     9.5
Name: Age, Length: 2923, dtype: float64
```

y_test

```
668     14.5
1580     9.5
3784    12.5
463      6.5
2615    13.5
...
1052    13.5
3439     9.5
1174    10.5
2210    10.5
2408    16.5
Name: Age, Length: 1254, dtype: float64
```

Question-11. Build the model

Solution:

```
from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
predict = lr.predict(x_test)
```

predict

```
array([12.32892845, 10.26905996, 11.99728864, ..., 11.36906433,
       13.83936664, 11.45100404])
```


Question-12. Train the model

Solution:

y_train

```
1376    11.5
1225     6.5
2722     8.5
3387    10.5
2773    12.5
...
1033    11.5
3264    13.5
1653    11.5
2607    10.5
2732     9.5
Name: Age, Length: 2923, dtype: float64
```

Question-13. Test the model

Solution:

y_test

```
668     14.5
1580     9.5
3784    12.5
463      6.5
2615    13.5
...
1052    13.5
3439     9.5
1174    10.5
2210    10.5
2408    16.5
Name: Age, Length: 1254, dtype: float64
```

Question-14. Measure the performance using Metrics

Solution:

```
from sklearn.metrics import r2_score, mean_squared_error
mse = mean_squared_error(y_test, predict)
rmse = np.sqrt(mse)
print("mse = ", mse)
print("rmse = ", rmse)
r2_score(y_test, predict)
```

```
mse = 2.9684942599827626
rmse = 1.7229318790894672
0.4607933491755676
```