

Abalone Age Prediction

Assignment -3

Assignment Date	03 October 2022
Team ID	PNT2022TMID27812
Project Name	Smart Lender-Application Credibility Prediction for loan Approval
Student Name	S.G.Mydhrayan
Student Roll Number	311519104036
Maximum Marks	2 Marks

Question-1.Download dataset

Solution:

1	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings		
2	0.365	0.095	0.514	0.2245	0.101	0.15	15		
3	0.265	0.09	0.2255	0.0995	0.0485	0.07	7		
4	0.42	0.135	0.677	0.2565	0.1415	0.21	9		
5	0.365	0.125	0.516	0.2155	0.114	0.155	10		
6	0.255	0.08	0.205	0.0895	0.0395	0.055	7		
7	0.3	0.095	0.3515	0.141	0.0775	0.12	8		
8	0.415	0.15	0.7775	0.237	0.1415	0.33	20		
9	0.425	0.125	0.768	0.294	0.1495	0.26	16		
10	0.37	0.125	0.5095	0.2165	0.1125	0.165	9		
11	0.44	0.15	0.8945	0.3145	0.151	0.32	19		
12	0.38	0.14	0.6065	0.194	0.1475	0.21	14		
13	0.35	0.11	0.406	0.1675	0.081	0.135	10		
14	0.38	0.135	0.5415	0.2175	0.095	0.19	11		
15	0.405	0.145	0.6845	0.2725	0.171	0.205	10		
16	0.355	0.1	0.4755	0.1675	0.0805	0.185	10		
17	0.4	0.13	0.6645	0.258	0.133	0.24	12		
18	0.28	0.085	0.2905	0.095	0.0395	0.115	7		
19	0.34	0.1	0.451	0.188	0.087	0.13	10		
20	0.295	0.08	0.2555	0.097	0.043	0.1	7		
21	0.32	0.1	0.381	0.1705	0.075	0.115	9		
22	0.28	0.095	0.2455	0.0955	0.062	0.075	11		
23	0.275	0.1	0.2255	0.08	0.049	0.085	10		
24	0.44	0.155	0.9395	0.4275	0.214	0.27	12		
25	0.415	0.135	0.7635	0.318	0.21	0.2	9		

Question-2. Load the dataset

Solution:

```
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt
import sklearn

data = pd.read_csv(r"abalone.csv")
data.head
```

```
<bound method NDFrame.head of
0      M  0.455  0.365  0.095  0.5140  0.2245
1      M  0.350  0.265  0.090  0.2255  0.0995
2      F  0.530  0.420  0.135  0.6770  0.2565
3      M  0.440  0.365  0.125  0.5160  0.2155
4      I  0.330  0.255  0.080  0.2050  0.0895
...
4172   F  0.565  0.450  0.165  0.8870  0.3700
4173   M  0.590  0.440  0.135  0.9660  0.4390
4174   M  0.600  0.475  0.205  1.1760  0.5255
4175   F  0.625  0.485  0.150  1.0945  0.5310
4176   M  0.710  0.555  0.195  1.9485  0.9455

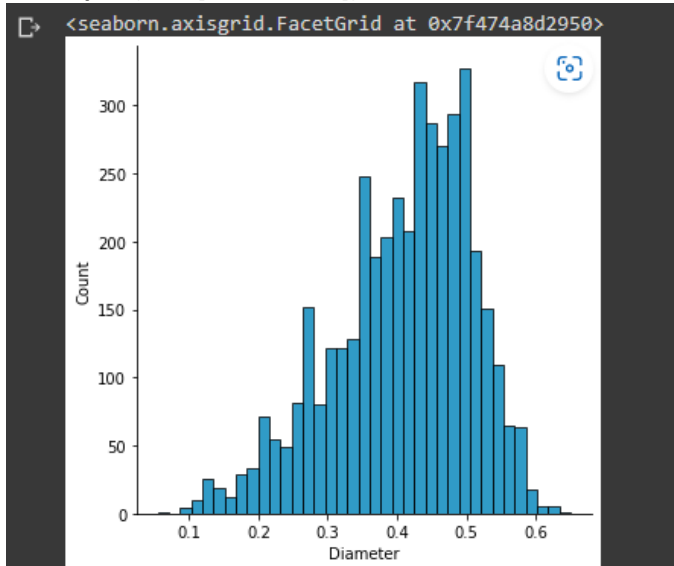
Viscera weight  Shell weight  Rings
0              0.1010      0.1500    15
1              0.0485      0.0700     7
2              0.1415      0.2100     9
3              0.1140      0.1550    10
4              0.0395      0.0550     7
...
4172           0.2390      0.2490    11
4173           0.2145      0.2605    10
4174           0.2875      0.3080     9
4175           0.2610      0.2960    10
4176           0.3765      0.4950    12
```

Question-3. Perform Below Visualizations.


3.1 Univariate Analysis

Solution:

`sns.displot(data[' Diameter'])`

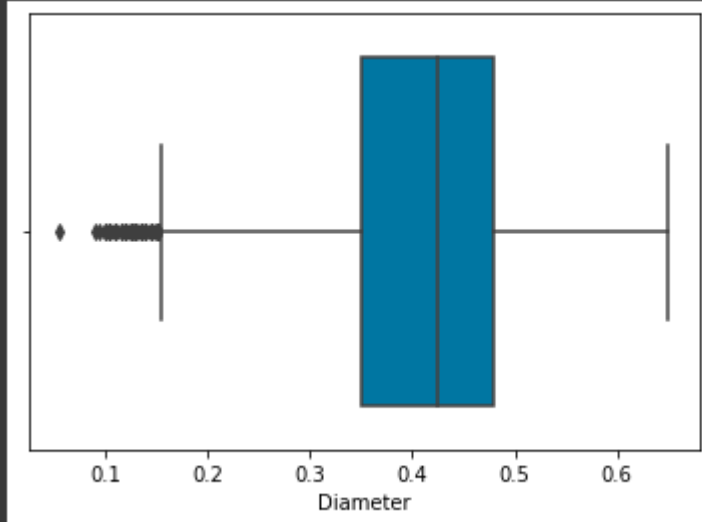


`sns.histplot(data[' Diameter '])`



```
sns.boxplot(x = data[' Diameter '])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f4731587050>
```




3.2 Bivariate Analysis

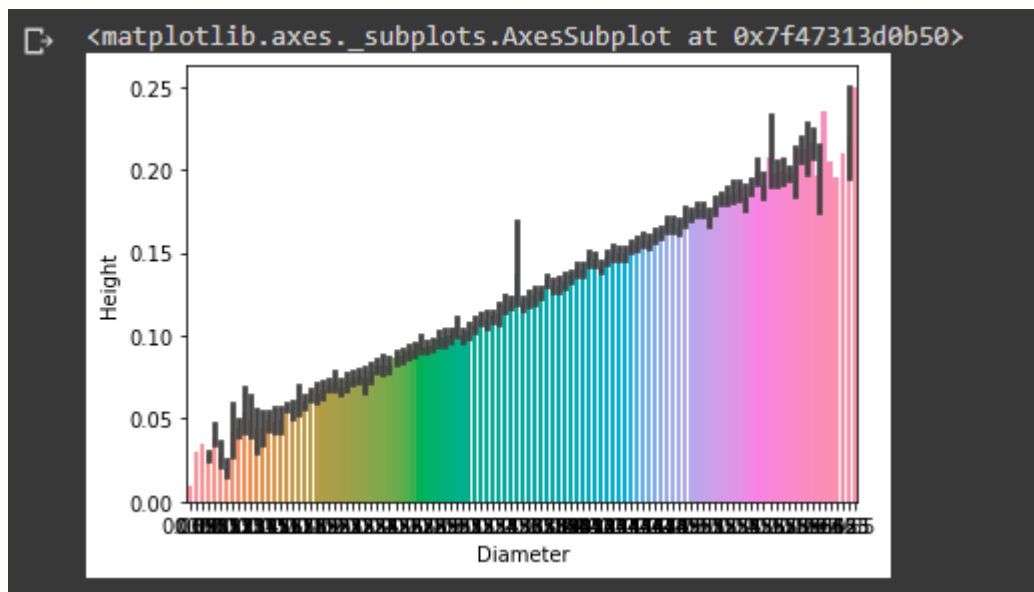
Solution:

```
sns.lineplot(data = data, x = 'Diameter', y = 'Rings')
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f47314ec350>
```




```
sns.barplot(data = data, x = 'Diameter',y = 'Height')
```



```
plt.figure(figsize=(20,20))
```


```
sns.barplot(data = data, x = 'Diameter', y = 'Shell weight')
```




3.3 Multivariate Analysis

Solution:

```
sns.scatterplot(data = data, x = 'Shell weight', y = 'Rings', hue = 'Length')
```



sns.pairplot(data)



Question-4. Perform descriptive statistics on the dataset.

Solution:

`data.mean(numeric_only = True)`

```
Length      0.523992
Diameter    0.407881
Height      0.139516
Whole weight 0.828742
Shucked weight 0.359367
Viscera weight 0.180594
Shell weight 0.238831
Rings       9.933684
dtype: float64
```

`data.median(numeric_only = True)`

```
Length      0.5450
Diameter    0.4250
Height      0.1400
Whole weight 0.7995
Shucked weight 0.3360
Viscera weight 0.1710
Shell weight 0.2340
Rings       9.0000
dtype: float64
```

`data['Whole weight'].mode()`

```
0 0.2225 dtype: float64
```

`data['Length'].mode()`

```
0 0.550 1 0.625 dtype: float64
```

`data['Rings'].unique()`

```
array([15, 7, 9, 10, 8, 20, 16, 19, 14, 11, 12, 18, 13, 5, 4, 6, 21,
       17, 22, 1, 3, 26, 23, 29, 2, 27, 25, 24])
```

`data.std(numeric_only=True)`

```
Length      0.120093
Diameter    0.099240
Height      0.041827
Whole weight 0.490389
Shucked weight 0.221963
Viscera weight 0.109614
Shell weight 0.139203
Rings       3.224169
dtype: float64
```



```
data.describe()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.407881	0.139516	0.828742	0.359367	0.180594	0.238831	9.933684
std	0.120093	0.099240	0.041827	0.490389	0.221963	0.109614	0.139203	3.224169
min	0.075000	0.055000	0.000000	0.002000	0.001000	0.000500	0.001500	1.000000
25%	0.450000	0.350000	0.115000	0.441500	0.186000	0.093500	0.130000	8.000000
50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000	0.234000	9.000000
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000	0.329000	11.000000
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000	1.005000	29.000000

```
data['Whole weight'].value_counts()
```

```
0.2225    8
1.1345    7
0.9700    7
0.4775    7
0.1960    7
..
0.0475    1
1.8930    1
1.8725    1
2.1055    1
1.9485    1
Name: Whole weight, Length: 2429, dtype: int64
```

Question-5.Handle the Missing values.

Solution:

data.isnull().any()

```
Sex          False
Length       False
Diameter     False
Height       False
Whole weight False
Shucked weight False
Viscera weight False
Shell weight False
Rings        False
dtype: bool
```


data.isnull().sum()

```
Sex          0
Length       0
Diameter     0
Height       0
Whole weight 0
Shucked weight 0
Viscera weight 0
Shell weight 0
Rings        0
dtype: int64
```

Question-6. Find the outliers and replace the outliers

Solution:

```
sns.boxplot(x = data['Rings'])
```



```
q = data.quantile([0.75,0.25])
```

```
q
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings
0.75	0.615	0.48	0.165	1.1530	0.502	0.2530	0.329	11.0
0.25	0.450	0.35	0.115	0.4415	0.186	0.0935	0.130	8.0

```
iqr = q.iloc[0] - q.iloc[1]
```

```
iqr
```

```
Length      0.1650
Diameter    0.1300
Height      0.0500
Whole weight 0.7115
Shucked weight 0.3160
Viscera weight 0.1595
Shell weight 0.1990
Rings       3.0000
dtype: float64
```

```
u = q.iloc[0] + (1.5*iqr)
```

```
u
```

```
Length      0.86250
Diameter    0.67500
Height      0.24000
Whole weight 2.22025
Shucked weight 0.97600
Viscera weight 0.49225
Shell weight 0.62750
Rings       15.50000
dtype: float64
```


```
l = q.iloc[1] - (1.5*iqr)
```

```
l
```

```
Length      0.20250
Diameter    0.15500
Height      0.04000
Whole weight -0.62575
Shucked weight -0.28800
Viscera weight -0.14575
Shell weight -0.16850
Rings        3.50000
dtype: float64
```

```
data['Rings'] = np.where(np.logical_or(data['Rings']>15, data['Rings']<4), 9, data['Rings'])
sns.boxplot(x = data['Rings'])
```

```
<matplotlib.axes._subplots.AxesSubplot at 0x7f472d26c310>
```



Question-7. Check for Categorical columns and perform encoding

Solution:

```
data['Age'] = data['Rings'] + 1.5  
data['Sex'].value_counts()
```

```
M    1528  
I    1342  
F    1307  
Name: Sex, dtype: int64
```

```
from sklearn.preprocessing import LabelEncoder, OneHotEncoder  
le = LabelEncoder()  
data = pd.get_dummies(data)  
data.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Rings	Age	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	15	16.5	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	7	8.5	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	9	10.5	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	10	11.5	0	0	1
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	7	8.5	0	1	0

Question-8. Split the data into dependent and independent variables split the data in X and Y

Solution:

```
y = data['Age']  
data = data.drop(['Rings','Age'], axis=1)  
x = data  
data.head
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	Sex_F	Sex_I	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	0	0	1
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	0	0	1
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	1	0	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	0	0	1
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	0	1	0

y

```
0      16.5
1       8.5
2      10.5
3      11.5
4       8.5
...
4172    12.5
4173    11.5
4174    10.5
4175    11.5
4176    13.5
Name: Age, Length: 4177, dtype: float64
```

Question-9.Scale the independent variables

Solution:

```
from sklearn.preprocessing import StandardScaler, MinMaxScaler
sc = StandardScaler()
x_scaled = sc.fit_transform(x)
x_scaled
```

```
array([[ -0.57455813, -0.43214879, -1.06442415, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-1.44898585, -1.439929  , -1.18397831, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.05003309,  0.12213032, -0.10799087, ...,  1.48184628,
        -0.68801788, -0.75948762],
       ...,
       [ 0.6329849  ,  0.67640943,  1.56576738, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.84118198,  0.77718745,  0.25067161, ...,  1.48184628,
        -0.68801788, -0.75948762],
       [ 1.54905203,  1.48263359,  1.32665906, ..., -0.67483383,
        -0.68801788,  1.31667716]])
```

Question-10.Split data into Training and Testing

Solution:

```
from sklearn.model_selection import train_test_split
```

```
x_train, x_test, y_train, y_test = train_test_split(x_scaled, y, test_size = 0.3, random_state = 0)
```

x_train

```
array([[ 0.79954256,  1.0291325 ,  0.84844242, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-1.49062526, -1.54070702, -1.30353247, ..., -0.67483383,
         1.45345059, -0.75948762],
       [-1.24078877, -1.33915098, -1.06442415, ..., -0.67483383,
         1.45345059, -0.75948762],
       ...,
       [ 0.59134549,  0.42446438,  0.13111745, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.84118198,  0.82757646,  0.6093341 , ...,  1.48184628,
        -0.68801788, -0.75948762],
       [-0.94931287, -0.83526087, -0.70576167, ..., -0.67483383,
         1.45345059, -0.75948762]])
```

x_train.shape

```
(2923, 10)
```

x_test

```
array([[ 0.21659075,  0.17251933,  0.37022577, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [-0.1998034 , -0.07942572, -0.46665335, ..., -0.67483383,
         1.45345059, -0.75948762],
       [ 0.79954256,  0.72679844,  0.37022577, ..., -0.67483383,
        -0.68801788,  1.31667716],
       ...,
       [ 0.92446081,  0.87796547, -2.97729071, ...,  1.48184628,
        -0.68801788, -0.75948762],
       [ 1.13265788,  0.97874349,  1.44621322, ..., -0.67483383,
        -0.68801788,  1.31667716],
       [ 0.79954256,  0.77718745,  0.72888826, ..., -0.67483383,
         1.45345059, -0.75948762]])
```

x_test.shape

```
(1254, 10)
```

y_train

```
1376    11.5
1225     6.5
2722     8.5
3387    10.5
2773    12.5
...
1033    11.5
3264    13.5
1653    11.5
2607    10.5
2732     9.5
Name: Age, Length: 2923, dtype: float64
```

y_test

```
668     14.5
1580     9.5
3784    12.5
463      6.5
2615    13.5
...
1052    13.5
3439     9.5
1174    10.5
2210    10.5
2408    16.5
Name: Age, Length: 1254, dtype: float64
```

Question-11. Build the model

Solution:

```
from sklearn.linear_model import LinearRegression
```

```
lr = LinearRegression()
```

```
lr.fit(x_train,y_train)
```

```
LinearRegression()
```

```
predict = lr.predict(x_test)
```

```
predict
```

```
array([12.32892845, 10.26905996, 11.99728864, ..., 11.36906433,
       13.83936664, 11.45100404])
```


Question-12. Train the model

Solution:

y_train

```
1376    11.5
1225     6.5
2722     8.5
3387    10.5
2773    12.5
...
1033    11.5
3264    13.5
1653    11.5
2607    10.5
2732     9.5
Name: Age, Length: 2923, dtype: float64
```

Question-13. Test the model

Solution:

y_test

```
668     14.5
1580     9.5
3784    12.5
463      6.5
2615    13.5
...
1052    13.5
3439     9.5
1174    10.5
2210    10.5
2408    16.5
Name: Age, Length: 1254, dtype: float64
```

Question-14. Measure the performance using Metrics

Solution:

```
from sklearn.metrics import r2_score, mean_squared_error
mse = mean_squared_error(y_test, predict)
rmse = np.sqrt(mse)
print("mse = ", mse)
print("rmse = ", rmse)
r2_score(y_test, predict)
```

```
mse = 2.9684942599827626
rmse = 1.7229318790894672
0.4607933491755676
```